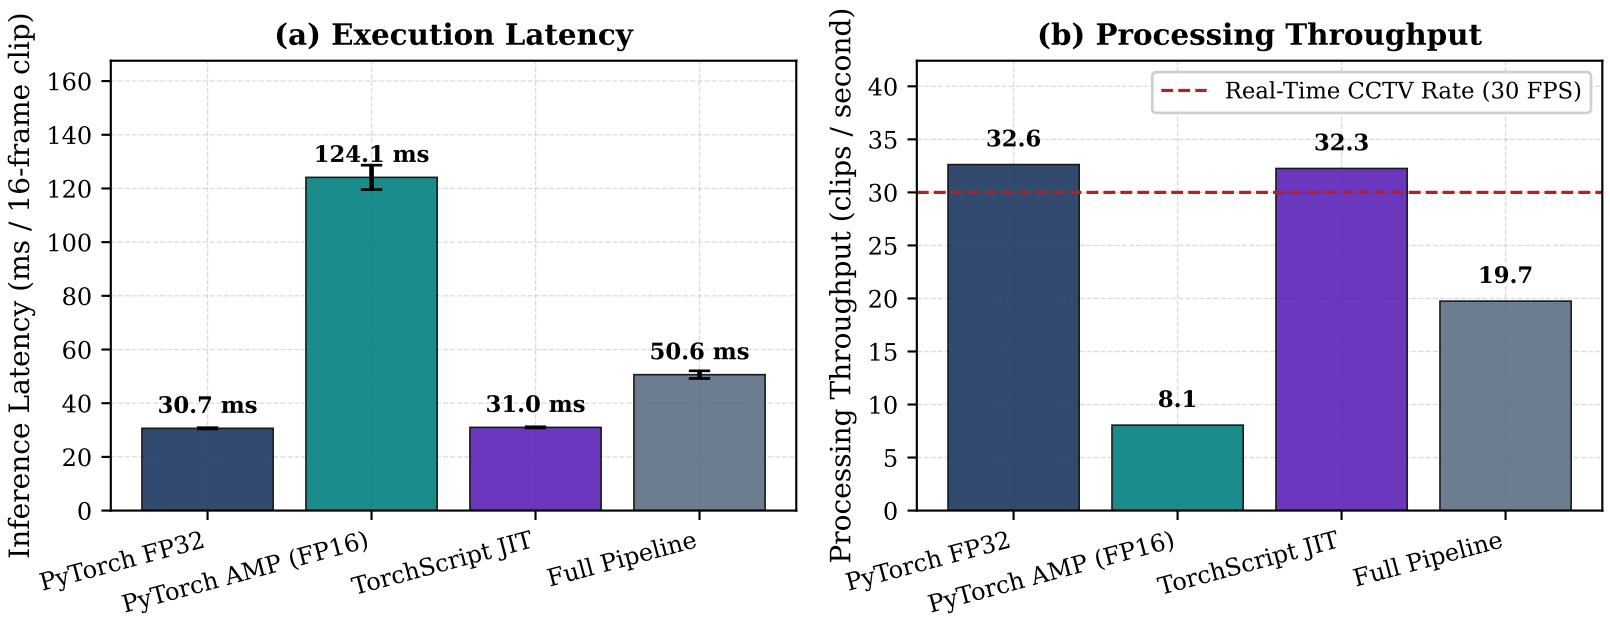
(a) Execution Latency
Inference Latency (ms / 16-frame clip)
0
20
40
60
80
100
120
140
160
30.7 ms
124.1 ms
31.0 ms
50.6 ms
PyTorch FP32
PyTorch AMP (FP16)
TorchScript JIT
Full Pipeline
(b) Processing Throughput
Processing Throughput (clips / second)
0
5
10
15
20
25
30
35
40
Real-Time CCTV Rate (30 FPS)
32.6
8.1
32.3
19.7
PyTorch FP32
PyTorch AMP (FP16)
TorchScript JIT
Full Pipeline